

Stock price prediction

Dataset


- Relatively large stock market dataset from Kaggle
- Full historical daily price and volume data for all US-based stocks and ETFs trading on the NYSE, NASDAQ, and NYSE MKT
- After preprocessing: 4571 stocks, 1155 ETFs
- Only a subset of the dataset is used due to computational limitations
- Date, Open, High, Low, Close, Volume (daily)

EDA



- Part of the data was empty or corrupted
- We only kept files with at least 300 entries (\approx 10 months)
- Used a correlation matrix and PCA + K-means
- ETFs were much more correlated (as expected)

Example from EDA

- Candlestick plot of BYM with additional technical indicators





EDA of ETFs



Backtesting

- We backtested common trading strategies to serve as a benchmark for comparison
- EMA, MACD, Bollinger Bands, RSI and the simple buy & hold strategies were tested (with several parameters)
- Markets are hard to beat: buy & hold performed fairly well on most stocks and ETFs

Examples from backtesting



MDP for trading

- $\text{MDP} = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$
- $s_t = [p_t, p_{t-1}, \dots, p_{t-k}]$ is a vector of closing prices k units back in time
- $\mathcal{A} = \{\text{Buy, Hold, Sell}\}$
- $r_t = w_{t+1} - w_t$ means the reward at time t is the portfolio value at time (t + 1) minus the value at time t
- $Q(s, a) = \mathbb{E} \left[r + \gamma \max_{a'} Q(s', a') \mid s, a \right]$
- The former equation defines the expected return from taking action a in state s and then following the optimal policy

DQN

- When the number of states is large we approximate the Q-value with a neural network
- It's basically a neural network that takes a state as input and outputs a score (Q-value) for each action (Buy, Sell, Hold)
- Experience Replay: Instead of learning only from the latest trade, the agent remembers past trades and learns from random ones
- This prevents it from overreacting to recent events

PPO

- Policy gradient algorithm, meaning it directly learns the policy — a function that maps states to actions — by optimizing expected reward
- It learns a stochastic policy
- Policy: $\pi_{\theta}(a \mid s)$ is the probability that we take action a in state s (parameterized by a neural network)
- Advantage function: $A_t = Q(s_t, a_t) - V(s_t)$ tells the agent how much better (or worse) an action was compared to the average
- The goal is to maximize expected advantage while clipping changes that are too large

DDPG

- Designed for problems where actions are continuous, not discrete
- In trading, instead of just buy / hold / sell, we can decide how much to buy or sell — like allocating 23% of your capital to an asset
- Actor-Critic Framework: Actor is the decision maker, Critic is the evaluation function

DDPG: continued

- Observe market state s_t
- Actor outputs action $a_t = \mu(s_t)$
- Execute action: get reward r_t and next state s_{t+1}
- Critic learns how good that action was
- Actor updates to make better decisions next time

RL vs Backtest

